

# Homework - 3

Group 04

## Problem 1: Gradient Descent Algorithm for Multiple Linear Regression

The file concrete.csv includes 1,030 types of concrete with numerical features indicating characteristics of the concrete. The variable “strength” is treated as the response variable.

- Standardize all variables (including the response variable “strength”). Split the data set into a training set (60%) and a validation set (40%).
- Implement the gradient descent algorithm in R with the ordinary least square cost function.
- Fit the multiple linear regression model using the gradient descent algorithm and the training set. Try out different learning rates:  $\alpha = 0.01, 0.1, 0.3, 0.5$  and compare the speed of convergence by plotting the cost function. Determine the number of iterations needed for each  $\alpha$  value.
- Apply the fitted regression model to the validation set and evaluate the model performance (ME, RMSE, MAE, MPE, MPAAE). Calculate the correlation between the predicted strength and the actual strength. Create a lift chart to show model performance.

```
# Import Required Packages
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(data.table)
```

```
##
```

```
## Attaching package: 'data.table'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      between, first, last
```

```
library(reshape2)
```

```
##
```

```
## Attaching package: 'reshape2'
```

```
## The following objects are masked from 'package:data.table':
```

```
##
```

```
##      dcast, melt
library(MLmetrics)

##
## Attaching package: 'MLmetrics'
## The following object is masked from 'package:base':
##
##      Recall
library(ggplot2)

# Read the csv file
df <- data.table(read.csv("concrete.csv"))

# Scale the dataframe
df <- as.data.frame(scale(df))

# Split into train and validation datasets
training_rows <- sample(seq_len(nrow(df)), size = floor(0.6 * nrow(df)))

train_data <- df[training_rows, ]
validation_data <- df[-training_rows, ]
```

Implementing Gradient Descent algorithm with the Ordinary Least Square cost function.

```
# Define the gradient descent function
gradient_desc <- function(x, y, lr, iters) {
  # First we create a list to keep the track
  # of the cost function for each iteration
  losses <- list()

  # Convert y to a matrix
  y <- as.matrix(y)

  # create a column of 1
  ones <- rep(1, dim(x)[[1]])
  # append it to the input (this is our X0)
  X <- as.matrix(cbind(ones, x))
  # Calculate number of samples
  n <- length(y)

  # Initialize model parameters/coefficients
  theta <- as.matrix(rnorm(n = dim(X)[2], 0, 1))

  # Calculate model predictions
  y_hat <- X %*% theta

  # calculate the loss using OLS cost function
  loss <- sum((y_hat - y)^2) / (2 * n)

  # Calculate the gradients of the cost function
  grads <- t(X) %*% (y_hat - y)

  # Update theta
```

```

theta <- theta - lr * (1 / n) * grads

# That was the first iteration of the gradient descent algorithm
# Let's add the cost function to the list
losses[[1]] <- loss

counter <- 0
# Number of iterations required to get the lowest loss
sufficient_iterations <- 0
for (i in 1:iters) {
  # Calculate model predictions
  y_hat <- X %*% theta

  # Calculate the loss using OLS cost function
  loss <- sum((y_hat - y)^2) / (2 * n)

  # Calculate the gradients
  grads <- t(X) %*% (y_hat - y)

  # Update theta
  theta <- theta - lr * (1 / n) * grads

  # Add cost to the list
  losses[[i + 1]] <- loss

  if (round(losses[[i]], 4) <= round(loss, 4)) {
    if (counter > 6) {
      break
    } else {
      counter <- counter + 1
      sufficient_iterations <- sufficient_iterations + 1
    }
  } else {
    counter <- 0
    sufficient_iterations <- sufficient_iterations + 1
  }
}

sufficient_iterations <- sufficient_iterations - counter
# return the theta (aka model weights)
return(list(
  "coeffs" = theta,
  "losses" = losses,
  "iterations_required" = sufficient_iterations,
  "final_loss" = loss
))
}

# Predict function
predict <- function(x, theta) {

```

```

ones <- rep(1, dim(x)[[1]])
# append it to the input (this is our X0)
X <- as.matrix(cbind(ones, x))

return(X %*% t(theta))
}

```

Now we create and train 4 models each with a different learning rate

```

# Model 1, lr = 0.01
model1 <- gradient_desc(train_data[, 1:8], train_data$strength, lr = 0.01, iters = 10000)

model1_weights <- t(model1$coeffs)
model1_losses <- melt(data.frame(model1$losses))
model1_losses$index <- 1:dim(model1_losses)[[1]]

# Model 2, lr = 0.10
model2 <- gradient_desc(train_data[, 1:8], train_data$strength, lr = 0.10, iters = 10000)

model2_weights <- t(model2$coeffs)
model2_losses <- melt(data.frame(model2$losses))
model2_losses$index <- 1:dim(model2_losses)[[1]]

# Model 3, lr = 0.30
model3 <- gradient_desc(train_data[, 1:8], train_data$strength, lr = 0.30, iters = 10000)

model3_weights <- t(model3$coeffs)
model3_losses <- melt(data.frame(model3$losses))
model3_losses$index <- 1:dim(model3_losses)[[1]]

# Model 4, lr = 0.50
model4 <- gradient_desc(train_data[, 1:8], train_data$strength, lr = 0.50, iters = 10000)

model4_weights <- t(model4$coeffs)
model4_losses <- melt(data.frame(model4$losses))
model4_losses$index <- 1:dim(model4_losses)[[1]]

```

Let's plot the loss vs number of iterations for each model to evaluate their performance.

```

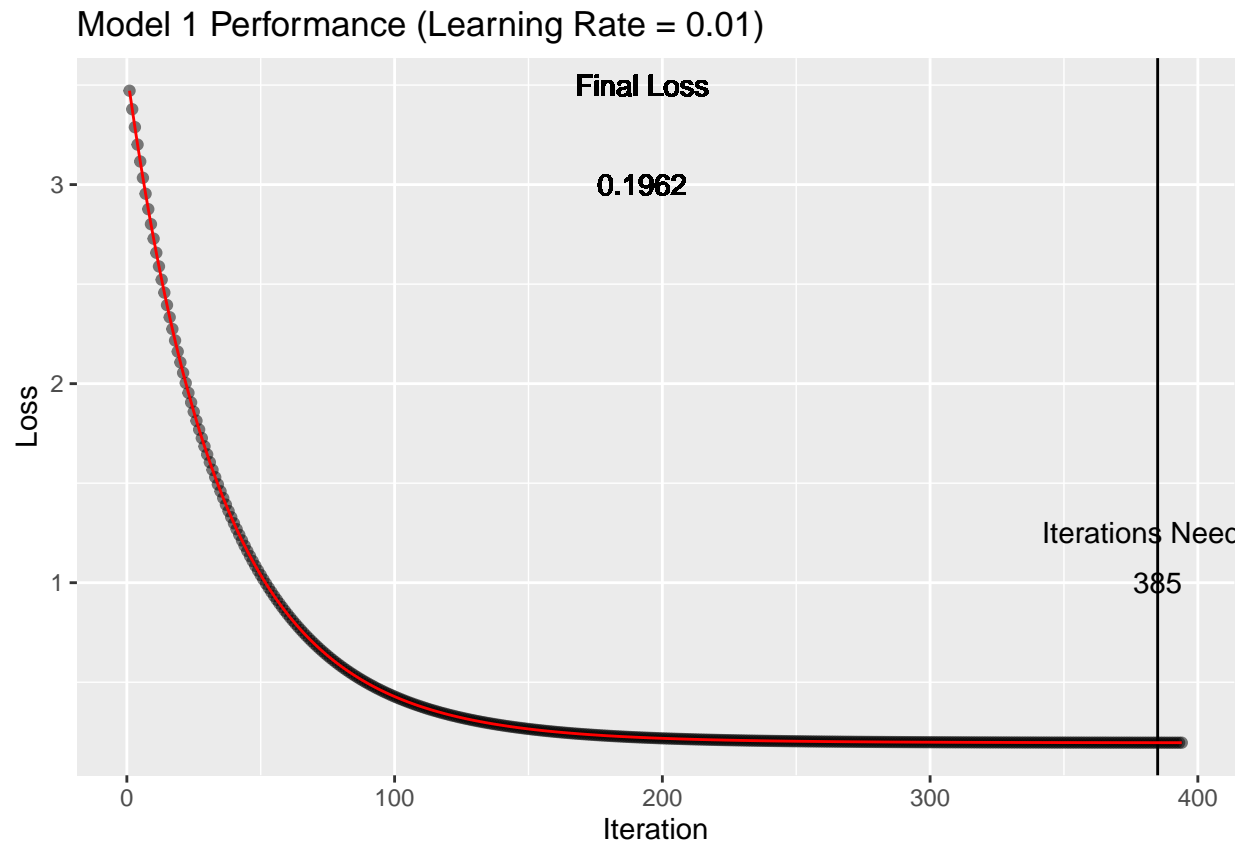
# Model 1
ggplot(model1_losses, aes(x = index, y = value)) +
  geom_point(alpha = 0.5) +
  geom_vline(xintercept = model1$iterations_required) +
  geom_text(x = model1$iterations_required / 2, y = 3.5, label = "Final Loss") +
  geom_text(x = model1$iterations_required / 2, y = 3, label = as.character(round(model1$final_loss, 4))) +
  geom_text(
    x = model1$iterations_required,
    y = 1,
    label = as.character(model1$iterations_required),
    check_overlap = TRUE
  ) +
  geom_text(
    x = model1$iterations_required,
    y = 1.25,

```

```

    label = "Iterations Needed",
    check_overlap = TRUE
  ) +
  geom_line(color = "red") +
  labs(x = "Iteration", y = "Loss") +
  ggtitle("Model 1 Performance (Learning Rate = 0.01)")

```




```

# Model 2
ggplot(model2_losses, aes(x = index, y = value)) +
  geom_point(alpha = 0.5) +
  geom_vline(xintercept = model2$iterations_required) +
  geom_text(x = model2$iterations_required / 2, y = 3.5, label = "Final Loss") +
  geom_text(x = model2$iterations_required / 2, y = 3, label = as.character(round(model2$final_loss, 4))) +
  geom_text(
    x = model2$iterations_required,
    y = 1,
    label = as.character(model2$iterations_required),
    check_overlap = TRUE
  ) +
  geom_text(
    x = model2$iterations_required,
    y = 1.25,
    label = "Iterations Needed",
    check_overlap = TRUE
  ) +
  geom_line(color = "red") +

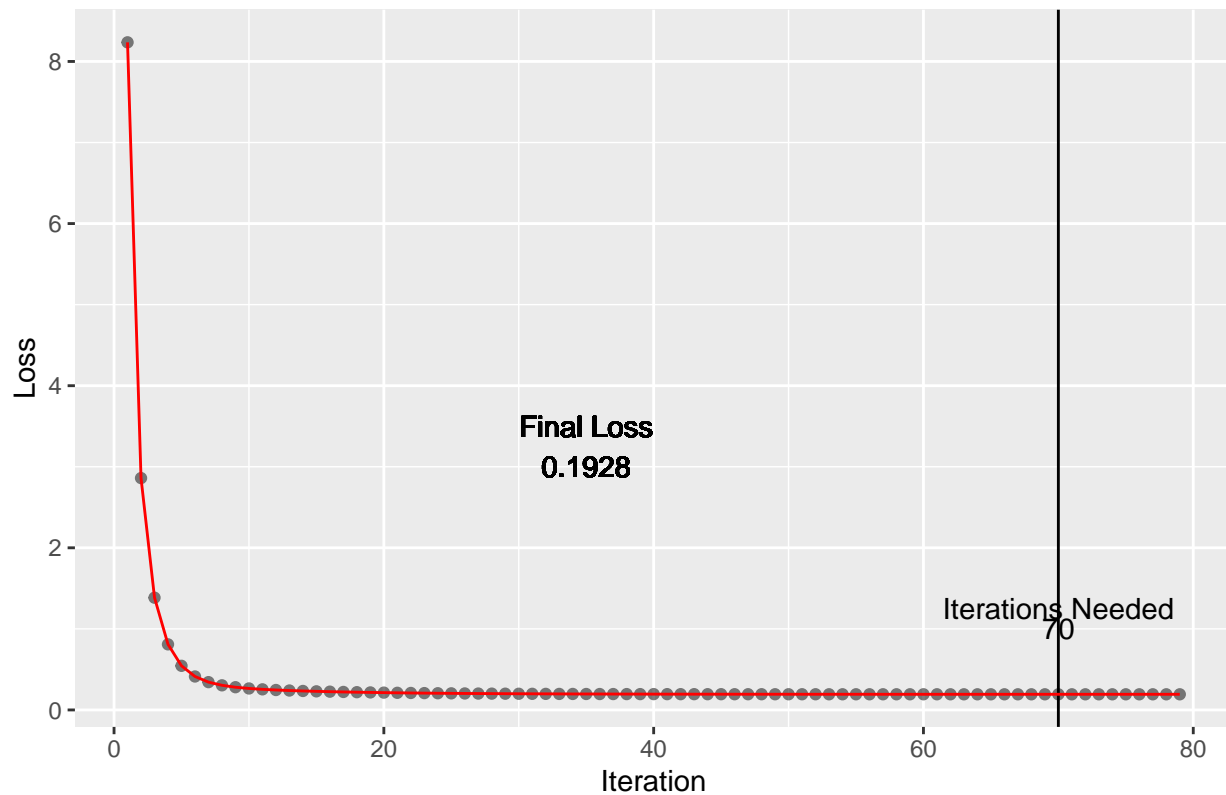
```

```
labs(x = "Iteration", y = "Loss") +
ggtitle("Model 2 Performance (Learning Rate = 0.10)")
```




```
# Model 3
ggplot(model3_losses, aes(x = index, y = value)) +
  geom_point(alpha = 0.5) +
  geom_vline(xintercept = model3$iterations_required) +
  geom_text(x = model3$iterations_required / 2, y = 3.5, label = "Final Loss") +
  geom_text(x = model3$iterations_required / 2, y = 3, label = as.character(round(model3$final_loss, 4))) +
  geom_text(
    x = model3$iterations_required,
    y = 1,
    label = as.character(model3$iterations_required),
    check_overlap = TRUE
  ) +
  geom_text(
    x = model3$iterations_required,
    y = 1.25,
    label = "Iterations Needed",
    check_overlap = TRUE
  ) +
  geom_line(color = "red") +
  labs(x = "Iteration", y = "Loss") +
  ggtitle("Model 3 Performance (Learning Rate = 0.30)")
```

Model 3 Performance (Learning Rate = 0.30)



```
# Model 4
ggplot(model4_losses, aes(x = index, y = value)) +
  geom_point(alpha = 0.5) +
  geom_vline(xintercept = model4$iterations_required) +
  geom_text(x = model4$iterations_required / 2, y = 3.5, label = "Final Loss") +
  geom_text(x = model4$iterations_required / 2, y = 3, label = as.character(round(model4$final_loss, 4))) +
  geom_text(
    x = model4$iterations_required,
    y = 1,
    label = as.character(model4$iterations_required),
    check_overlap = TRUE
  ) +
  geom_text(
    x = model4$iterations_required,
    y = 1.25,
    label = "Iterations Needed",
    check_overlap = TRUE
  ) +
  geom_line(color = "red") +
  labs(x = "Iteration", y = "Loss") +
  ggtitle("Model 4 Performance (Learning Rate = 0.50)")
```

Model 4 Performance (Learning Rate = 0.50)



```
cat("Number of iterations required for each model are :\n")
```

```
## Number of iterations required for each model are :
```

```
cat("Model 1:", as.character(model1$iterations_required), "\n")
```

```
## Model 1: 385
```

```
cat("Model 2:", as.character(model2$iterations_required), "\n")
```

```
## Model 2: 334
```

```
cat("Model 3:", as.character(model3$iterations_required), "\n")
```

```
## Model 3: 70
```

```
cat("Model 4:", as.character(model4$iterations_required), "\n")
```

```
## Model 4: 36
```

As observed, the model converges faster as the learning rate increases.

Testing the model on the validation data and calculating errors -

```
# We define the Mean Error function
```

```
ME <- function(y_hat, y) {  
  sum(y - y_hat) / length(y)  
}
```

```
# We define the Mean Percentage Error Function
```

```
MPE <- function(y_hat, y) {
```



```
(sum((y - y_hat) / y)) / length(y)
}
```

Now let's look at the model statistics -

```
model1_predictions <- predict(validation_data[, 1:8], model1_weights)
```

```
cat("----Model 1 Summary ----\n")
```

```
## ----Model 1 Summary ----
```

```
cat("MAE:", MAE(model1_predictions, validation_data[, 9]), "\n")
```

```
## MAE: 0.4862348
```

```
cat("RMSE:", RMSE(model1_predictions, validation_data[, 9]), "\n")
```

```
## RMSE: 0.6193269
```

```
cat("ME:", ME(model1_predictions, validation_data[, 9]), "\n")
```

```
## ME: -0.03444517
```

```
cat("MPE:", MPE(model1_predictions, validation_data[, 9]), "\n")
```

```
## MPE: 0.4457131
```

```
cat("MPAE", MAPE(model1_predictions, validation_data[, 9]), "\n")
```

```
## MPAE 1.813482
```

```
model2_predictions <- predict(validation_data[, 1:8], model2_weights)
```

```
cat("----Model 2 Summary ---- \n")
```

```
## ----Model 2 Summary ----
```

```
cat("MAE:", MAE(model2_predictions, validation_data[, 9]), "\n")
```

```
## MAE: 0.498266
```

```
cat("RMSE:", RMSE(model2_predictions, validation_data[, 9]), "\n")
```

```
## RMSE: 0.6270598
```

```
cat("ME:", ME(model2_predictions, validation_data[, 9]), "\n")
```

```
## ME: -0.01841473
```

```
cat("MPE:", MPE(model2_predictions, validation_data[, 9]), "\n")
```

```
## MPE: 0.2837325
```

```
cat("MPAE", MAPE(model2_predictions, validation_data[, 9]), "\n")
```

```
## MPAE 2.111422
```

```
model3_predictions <- predict(validation_data[, 1:8], model3_weights)
```

```
cat("----Model 3 Summary ----\n")
```

```
## ----Model 3 Summary ----
```

```

cat("MAE:", MAE(model3_predictions, validation_data[, 9]), "\n")

## MAE: 0.4914922
cat("RMSE:", RMSE(model3_predictions, validation_data[, 9]), "\n")

## RMSE: 0.6208526
cat("ME:", ME(model3_predictions, validation_data[, 9]), "\n")

## ME: -0.01422963
cat("MPE:", MPE(model3_predictions, validation_data[, 9]), "\n")

## MPE: 0.3826415
cat("MPAE", MAPE(model3_predictions, validation_data[, 9]), "\n")

## MPAE 2.022136
model4_predictions <- predict(validation_data[, 1:8], model4_weights)

cat("----Model 4 Summary ----\n")

## ----Model 4 Summary ----
cat("MAE:", MAE(model4_predictions, validation_data[, 9]), "\n")

## MAE: 0.4916012
cat("RMSE:", RMSE(model4_predictions, validation_data[, 9]), "\n")

## RMSE: 0.6210676
cat("ME:", ME(model4_predictions, validation_data[, 9]), "\n")

## ME: -0.01566395
cat("MPE:", MPE(model4_predictions, validation_data[, 9]), "\n")

## MPE: 0.3318516
cat("MPAE", MAPE(model4_predictions, validation_data[, 9]), "\n")

## MPAE 2.004415

```

We can see that all the models have approximately the same accuracy regardless the learning rate.

Calculating the correlation between predicted strength and actual strength

```

cat("The correlation is :", cor(model1_predictions, validation_data[, 9]), "\n")

```

```

## The correlation is : 0.788954

```

Plotting a lift chart

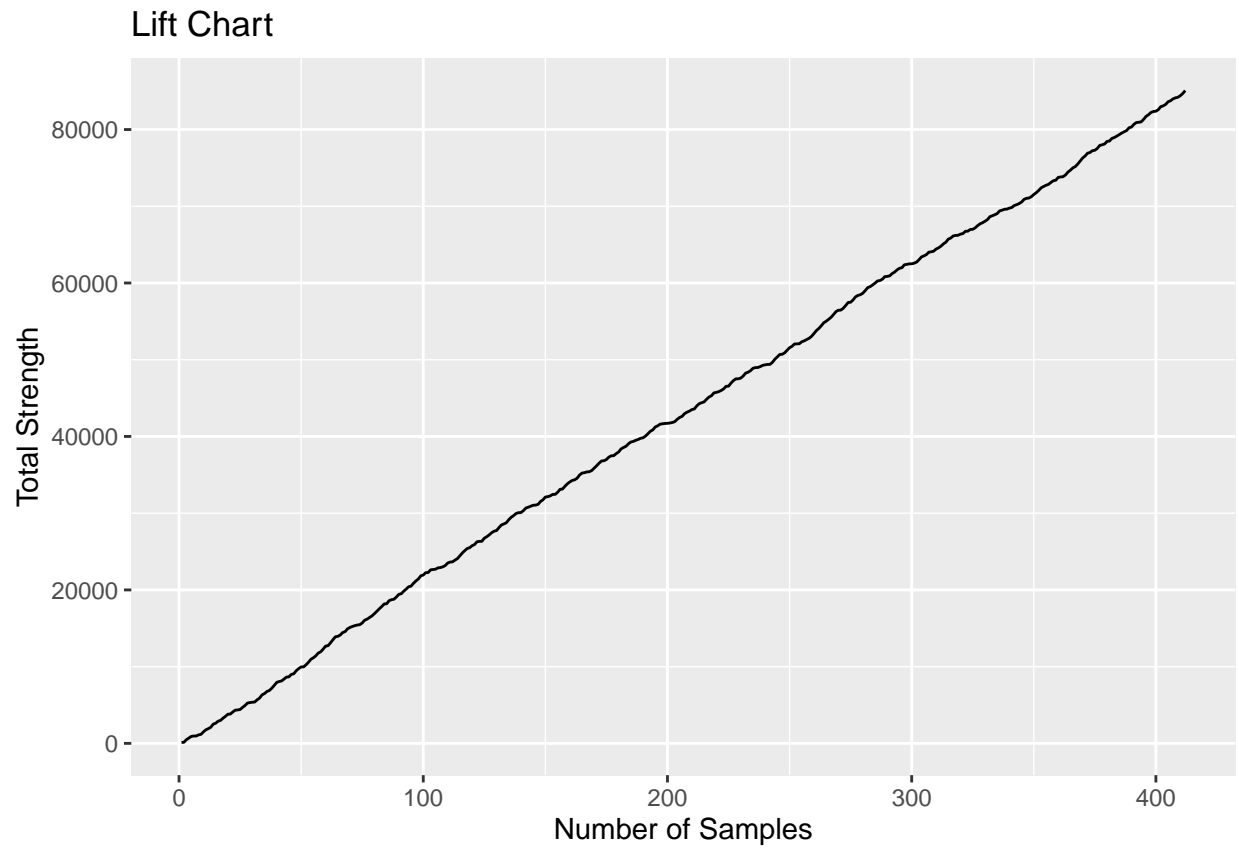
```

# Create a temp data frame to calculate the sumulative strength
temp <- data.frame("strength" = order(validation_data[, 9]))
temp$cumstrength <- cumsum(temp$strength)
temp$samples <- 1:dim(temp)[[1]]

# Plot the lift chart
ggplot(temp, aes(x = samples, y = cumstrength)) +
  geom_line() +

```

```
labs(x = "Number of Samples", y = "Total Strength") +  
ggtitle("Lift Chart")
```



```
# Delete all environment variables  
rm(list = ls())
```

---

## Problem 2